

Fraud Detection System

#import necessary models and libraries

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_predict
```

```
import collections
```

```
%matplotlib inline
```

```
import matplotlib.pyplot as plt
```

```
import data from kaggle: https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud
```

#read the data and save it in df

```
df = pd.read_csv("creditcard.csv")
df.head()
```

	Time	V1	V2	V3	V4	V5	V6	
V7 \								
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	
	0.239599							
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	
	0.078803							
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	
	0.791461							
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	
	0.237609							
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	
	0.592941							
		V8	V9	...	V21	V22	V23	V24
V25 \								
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	
	0.128539							
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	
	0.167170							
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	
	0.327642							
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	
	0.647376							
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	

0.206010

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

Preprocessing

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null float64
1   V1          284807 non-null float64
2   V2          284807 non-null float64
3   V3          284807 non-null float64
4   V4          284807 non-null float64
5   V5          284807 non-null float64
6   V6          284807 non-null float64
7   V7          284807 non-null float64
8   V8          284807 non-null float64
9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
```

```
30 Class    284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
# check whether any of the values are missing
df.isnull().sum()
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
```

```
# return the series containing counts of unique values
df['Class'].value_counts()
```

```
0    284315
1      492
Name: Class, dtype: int64
```

Split data into normal and fraud 0 - # of normal transactions 1 - # of fraud transactions

```
# seperate data for the further analysis
```

```
normal = df[df.Class == 0]
```

```
fraud = df[df.Class == 1]
```

```
# need to find out the matrices shapes
```

```
print(normal.shape)
```

```
print(fraud.shape)
```

```
(284315, 31)
```

```
(492, 31)
```

Need to view the statistics of both of the datasets

```
normal.Amount.describe()
```

```
count    284315.000000
```

```
mean      88.291022
```

```
std       250.105092
```

```
min        0.000000
```

```
25%        5.650000
```

```
50%       22.000000
```

```
75%       77.050000
```

```
max      25691.160000
```

```
Name: Amount, dtype: float64
```

```
fraud.Amount.describe()
```

```
count      492.000000
```

```
mean      122.211321
```

```
std       256.683288
```

```
min        0.000000
```

```
25%        1.000000
```

```
50%        9.250000
```

```
75%       105.890000
```

```
max      2125.870000
```

```
Name: Amount, dtype: float64
```

```
# compare the values for both datasets
```

```
df.groupby('Class').mean()
```

	Time	V1	V2	V3	V4	V5
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225

	V6	V7	V8	V9	...	V20	V21
...							

```

0      0.002419  0.009637 -0.000987  0.004467  ... -0.000644 -0.001235
1     -1.397737 -5.568731  0.570636 -2.581123  ...  0.372319  0.713588

```

```

          V22      V23      V24      V25      V26      V27
V28 \
Class
0     -0.000024  0.000070  0.000182 -0.000072 -0.000089 -0.000295 -
0.000131
1      0.014049 -0.040308 -0.105130  0.041449  0.051648  0.170575
0.075667

```

```

          Amount
Class
0      88.291022
1     122.211321

```

[2 rows x 30 columns]

Under Sampling

The undersampling is conducted to: make the set balanced avoid overfitting

```
# creating less samples to match normal dataset
```

```
normal_under_sample = normal.sample(n=492)
```

```
# build a new dataframe
```

```
normal_df = pd.concat([normal_under_sample, fraud], axis = 0)
```

```
normal_df.head()
```

```

          Time      V1      V2      V3      V4      V5
V6 \
95739    65442.0 -5.015087  3.041594  0.508471  2.479145 -1.703519
1.673319
118015    74915.0  1.345707 -0.575480  0.603205 -0.695274 -1.041123 -
0.501084
212093   138709.0  2.148813 -0.074116 -2.573260 -0.337782  1.000894 -
0.596484
116653    74374.0 -0.341830  1.184869  1.309382  0.055606  0.093200 -
0.964935
195058   130875.0  1.968937 -0.557725  0.030995  0.171990 -0.893319 -
0.015563

```

```

          V7      V8      V9  ...      V21      V22
V23 \

```

```

95739 -1.481565  1.522606  0.760693  ...  0.163183  0.150103
0.057565
118015 -0.745245 -0.023503 -0.656310  ...  0.135347  0.317288
0.044331
212093  0.502468 -0.201749  0.256163  ...  0.126097  0.365870 -
0.086476
116653  0.766500 -0.116748 -0.510403  ... -0.254564 -0.626968 -
0.022126
195058 -1.010437  0.196478  1.160642  ...  0.058546  0.307938
0.276465

```

```

          V24          V25          V26          V27          V28  Amount
Class
95739 -0.322835  0.005904  0.135251 -1.666894  0.723122   17.24
0
118015  0.088236  0.281741 -0.253498  0.030028  0.018793   16.00
0
212093 -0.105702  0.457918  0.247134 -0.104046 -0.094673    1.89
0
116653  0.354953 -0.134689  0.069284  0.250350  0.099980    3.99
0
195058 -0.311382 -0.585825  0.482921 -0.015300 -0.052749    1.00
0

```

[5 rows x 31 columns]

```

# need to check that we have the same counts for both datasets
normal_df['Class'].value_counts()

```

```

0    492
1    492
Name: Class, dtype: int64

```

```

# define X and Y
X = normal_df.drop(columns = 'Class', axis = 1)
Y = normal_df["Class"]

```

```

print(X)

```

```

          Time          V1          V2          V3          V4          V5
V6 \
95739  65442.0 -5.015087  3.041594  0.508471  2.479145 -1.703519
1.673319
118015  74915.0  1.345707 -0.575480  0.603205 -0.695274 -1.041123 -
0.501084
212093 138709.0  2.148813 -0.074116 -2.573260 -0.337782  1.000894 -
0.596484
116653  74374.0 -0.341830  1.184869  1.309382  0.055606  0.093200 -
0.964935
195058 130875.0  1.968937 -0.557725  0.030995  0.171990 -0.893319 -
0.015563

```

...
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487 -
2.010494						
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581 -
1.326536						
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541 -
0.003346						
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618 -
2.943548						
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147 -
0.096695						
	V7	V8	V9	...	V20	V21
V22 \						
95739	-1.481565	1.522606	0.760693	...	0.001156	0.163183
0.150103						
118015	-0.745245	-0.023503	-0.656310	...	0.043921	0.135347
0.317288						
212093	0.502468	-0.201749	0.256163	...	-0.302968	0.126097
0.365870						
116653	0.766500	-0.116748	-0.510403	...	0.164019	-0.254564 -
0.626968						
195058	-1.010437	0.196478	1.160642	...	-0.178481	0.058546
0.307938						
...
.						
279863	-0.882850	0.697211	-2.064945	...	1.252967	0.778584 -
0.319189						
280143	-1.413170	0.248525	-1.127396	...	0.226138	0.370612
0.028234						
280149	-2.234739	1.210158	-0.652250	...	0.247968	0.751826
0.834108						
281144	-2.208002	1.058733	-1.632333	...	0.306271	0.583276 -
0.269209						
281674	0.223050	-0.068384	0.577829	...	-0.017652	-0.164350 -
0.295135						
	V23	V24	V25	V26	V27	V28
Amount						
95739	0.057565	-0.322835	0.005904	0.135251	-1.666894	0.723122
17.24						
118015	0.044331	0.088236	0.281741	-0.253498	0.030028	0.018793
16.00						
212093	-0.086476	-0.105702	0.457918	0.247134	-0.104046	-0.094673
1.89						
116653	-0.022126	0.354953	-0.134689	0.069284	0.250350	0.099980
3.99						
195058	0.276465	-0.311382	-0.585825	0.482921	-0.015300	-0.052749
1.00						

```

...      ...      ...      ...      ...      ...      ...
...
279863  0.639419 -0.294885  0.537503  0.788395  0.292680  0.147968
390.00
280143 -0.145640 -0.081049  0.521875  0.739467  0.389152  0.186637
0.76
280149  0.190944  0.032070 -0.739695  0.471111  0.385107  0.194361
77.89
281144 -0.456108 -0.183659 -0.328168  0.606116  0.884876 -0.253700
245.00
281674 -0.072173 -0.450261  0.313267 -0.289617  0.002988 -0.015309
42.53

```

[984 rows x 30 columns]

```

# splitting data into training and testing
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size =
0.2, stratify=Y, random_state = 2)

# initializing the values
X_train = X_train.values
X_test = X_test.values
y_train = y_train.values
y_test = y_test.values

```

Analysis

using ML algorithms and models

```

# create dict - classifiers for 4 models
classifiers = {
    "LogisticRegression": LogisticRegression(),
    "KNearest": KNeighborsClassifier(),
    "DecisionTreeClassifier": DecisionTreeClassifier()
}

# undersampling
from sklearn.model_selection import cross_val_score

for key, classifier in classifiers.items():
    classifier.fit(X_train, y_train)
    training_score = cross_val_score(classifier, X_train, y_train,
cv=5)
    print("Classifiers: ", classifier.__class__.__name__, "has a
training score of", round(training_score.mean(), 2) * 100, "% accuracy
score")

/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/
_logistic.py:458: ConvergenceWarning: lbfgs failed to converge

```



```
(status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.  
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

Classifiers: LogisticRegression has a training score of 94.0 % accuracy score

Classifiers: KNeighborsClassifier has a training score of 66.0 % accuracy score

Classifiers: SVC has a training score of 56.000000000000001 % accuracy score

Classifiers: DecisionTreeClassifier has a training score of 91.0 % accuracy score

Investigating Hyperparameters using Grid Search

```
# Hyperparameters for Logistic Regression
```

```
lr_hyp = {"penalty": ['l1', 'l2'], 'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000]}
```

```
grid_lr = GridSearchCV(LogisticRegression(max_iter = 200), lr_hyp,)  
grid_lr.fit(X_train, y_train)
```

```
# best estimator for the parameters
```

```
lr_be = grid_lr.best_estimator_
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/model_selection/  
_validation.py:378: FitFailedWarning:
```

```
35 fits failed out of a total of 70.
```

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting `error_score='raise'`.

Below are more details about the failures:

```
-----
-----
35 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.8/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File
"/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.8/dist-packages/sklearn/model_selection/_search.py:953: UserWarning: One or more of the test scores are non-finite: [
nan 0.92126099      nan 0.93646698      nan 0.93012174
      nan 0.93899057      nan 0.93008949      nan 0.9313795
      nan 0.9313795 ]
    warnings.warn(
```

Hyperparameters for KNeighbors

```
kn_hyp = {"n_neighbors": list(range(2,5,1)), 'algorithm': ['auto',
'ball_tree', 'kd_tree', 'brute']}
```

```
grid_kn = GridSearchCV(KNeighborsClassifier(), kn_hyp)
grid_kn.fit(X_train, y_train)
```

best estimator for the parameters

```
kn_be = grid_kn.best_estimator_
```

Hyperparameters for DecisionTree Classifier

```
tree_hp = {"criterion": ["gini", "entropy"], "max_depth":
list(range(2,4,1)),
           "min_samples_leaf": list(range(5,7,1))}
```

```
grid_tree = GridSearchCV(DecisionTreeClassifier(), tree_hp)
grid_tree.fit(X_train, y_train)
```

best estimator for the parameters

```
tree_be = grid_tree.best_estimator_
```

```
# Calculating and outputting the cross-validation scores of each model
log_reg_score = cross_val_score(lr_be, X_train, y_train, cv=5)
print('Logistic Regression Cross-Validation Score: ',
round(log_reg_score.mean() * 100, 2).astype(str) + '%')
```

```
knears_score = cross_val_score(kn_be, X_train, y_train, cv=5)
print('Knears Neighbors Cross-Validation Score',
round(knears_score.mean() * 100, 2).astype(str) + '%')
```

```
tree_score = cross_val_score(tree_be, X_train, y_train, cv=5)
print('DecisionTree Classifier Cross-Validation Score',
round(tree_score.mean() * 100, 2).astype(str) + '%')
```

```
Logistic Regression Cross-Validation Score: 93.9%
Knears Neighbors Cross-Validation Score 65.81%
DecisionTree Classifier Cross-Validation Score 92.63%
```

This shows that the best ML classifier or model for the Fraud Detection system is Logistic Regression and this can be validated through the plot

```
# build a plot for the Logistic Regression using learnign_curve
from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import learning_curve

import matplotlib.pyplot as plt

def plot_learning_curve(classifier1, X, y, ylim=None, cv=1,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0,
5)):
    f, ax1 = plt.subplots( 1,1, figsize=(20,14), sharey=True)
    if ylim is not None:
        plt.ylim(*ylim)
    # Logic Regression Learning Curve
    train_sizes, train_scores, test_scores = learning_curve(
        classifier1, X, y, cv=cv, n_jobs=n_jobs,
train_sizes=train_sizes)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    ax1.plot(train_sizes, train_scores_mean, 'o-', color="#00FF00",
        label="Training score")
    ax1.plot(train_sizes, test_scores_mean, 'o-', color="#000000",
        label="Cross-validation score")
    ax1.set_title("Logistic Regression Learning Curve", fontsize=14)
    ax1.set_xlabel('Training size (m)')
    ax1.set_ylabel('Score')
    ax1.grid(True)
    ax1.legend(loc="best")
```

`return plt`

```
cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=42)
plot_learning_curve(lr_be, X_train, y_train, (0.87, 1.01), cv=cv,
n_jobs=1)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.
py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
<module 'matplotlib.pyplot' from '/usr/local/lib/python3.8/dist-
packages/matplotlib/pyplot.py'>
```

